Idea in bullets

Vasilis Gkolemis

June 2021

1 Intro

ALE plots are the best feature effect technique, because blablabla. However, they have three drawbacks; Firstly, the do not provide a measure of uncertainty about the effect they calculate (in other words). Secondly, they sometimes calculate different depending on the number of bins (and we do not have a measure to choose which one is closer to reality). Finally, for making accurate ALE plots, we sometimes need to use variable-size bins.

2 Can we trust the ALE plot?

In ALE, the user defines the bin-size, through one of the following hyperparameters:

- K, number of bins
- dx, bin-size
- `min_points_per_bin`

which have '1-1' relation:

$$dx \Leftrightarrow K \Leftrightarrow \texttt{min\_points\_per\_bin}$$

Therefore, we refer to them interchangeably.

Statement

Altering the number of bins (K), leads to different ALE plots. Depending on K, we receive quite different ALE without any indicator which one is the correct. In other words, we need a metric to evaluate each plot.

Example

We define the following model:

$$f(x_1,x_2) = x_1x_2 + \begin{cases} -5+0.3x_1 & ,0 \leq x_1 < 20 \\ 1+7(x_1-20) & ,20 \leq x_1 < 40 \\ 141-1.5(x_1-40) & ,40 \leq x_1 < 60 \\ 111+0(x_1-60) & ,60 \leq x_1 < 80 \\ 111-5(x_1-80) & ,80 \leq x_1 < 100 \end{cases} \tag{1}$$

where $x_1 \perp x_2$ and $x_2 \sim \mathcal{N}(\mu = 0, \sigma = 4)$. Therefore, the gradients wrt. x_1 are:

$$\frac{\partial f}{\partial x_1}(x_1) = x_2 + \begin{cases} 0.3 & ,0 \leq x_1 < 20 \\ 7 & ,20 \leq x_1 < 40 \\ 1.5 & ,40 \leq x_1 < 60 \\ 0 & ,60 \leq x_1 < 80 \\ -5 & ,80 \leq x_1 < 100 \end{cases} \tag{2}$$

where $x_2 \sim \mathcal{N}(\mu = 0, \sigma = 4)$. The ground truth ALE is:

$$f_{\texttt{ALE}}(x_s) = c + \begin{cases} -5+0.3x_s & ,0 \leq x_s < 20 \\ 1+7(x_s-20) & ,20 \leq x_s < 40 \\ 141-1.5(x_s-40) & ,40 \leq x_s < 60 \\ 111+0(x_s-60) & ,60 \leq x_s < 80 \\ 111-5(x_s-80) & ,80 \leq x_s < 100 \end{cases} \tag{3}$$

where c is a normalizing constant.

We generate N $= 100$ data points uniformly in the region $[0, 100]$, i.e., $\mathcal{D} \sim \mathcal{U}(0, 100)$. The produced feature effect plot is shown in figure 1.

We observe that we get different plots for $K = \{3, 5, 20, 100\}$, without information which one to trust. If we set a threshold on `min_points_per_bin`, then we discard plots (c) and (d) beacuse they violate the threshold. In this case, among (a) and (b), we cannot know which one to trust (We cannot trust both, since they provide different effects).

For $K = 3$, the available resolution is smaller than the one required, leading to an erroneous estimation. For $K = 5$, the feature effect matches the correct resolution. For $K = 20, 100$, the feature effect provides the general trend, but with noisy artifacts due to small number of points per bin (violating `min_points_per_bin`).

Proposal

We propose standard error as a metric for informing **to what extend we should trust the feature effect plot**. **Standard error shows the expected error in the computation of the feature effect plot.** However, there are some constraints we must notice:

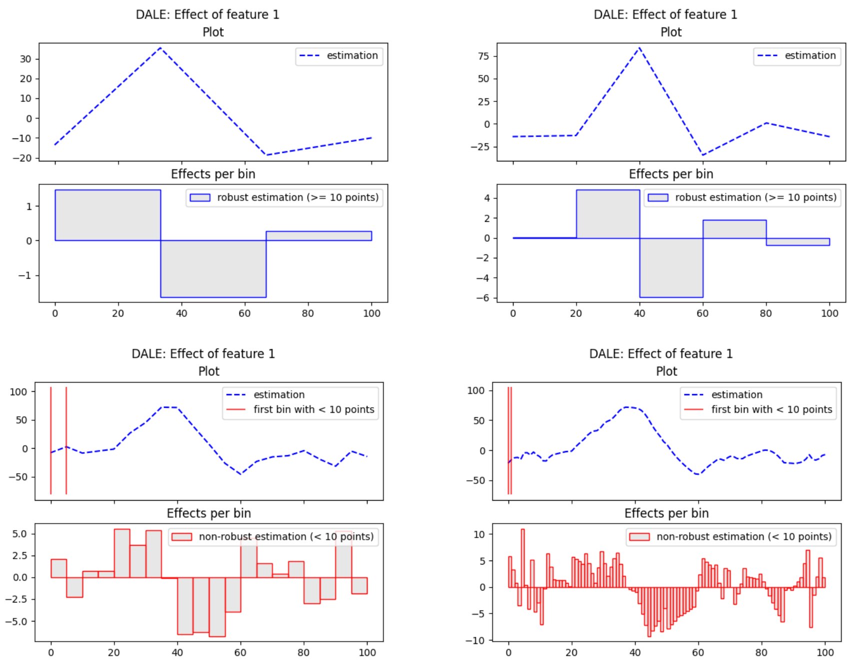

Figure 1: DALE effect for (a) $K = 3$, (b) $K = 5$, (c) $K = 20$, (d) $K = 100$

1. Our computations are based on the hypothesis that inside all bins **the gradient wrt. to the feature of interest doesn't depend on the value feature of interest**. In our example, in the interval $[0, 20)$ the gradient is $x_2 + 0.3$ (independent of x_2). But in the interval $[0, 40)$ the gradient is 0.3 if $x_1 < 20$ and 7 otherwise (not independent of x_2). Unfortunately, we cannot when this is the case. We just know that as the bins grow larger, it is more possible to violated this hypothesis, as in Figure 2(a). In this case both the ALE effect and the standard error are wrong.

2. The standard error should be trusted when it is estimated by a large of data points. For example, in plots (c) and (d), there are bins with less than 10 points. Therefore, in these cases, we cannot trust the plot or the standard error.

3. For being confident about the feature effect plot, we should check the region covered by 2 or 3 times the standard error. For example in 2, we observe that in plots (a) and (b), the green region showing the standard error is very big (covers the region from almost -100 to 100). Therefore, the effect cannot be trusted.

If the above criteria are met, the standard error is a useful metric. For

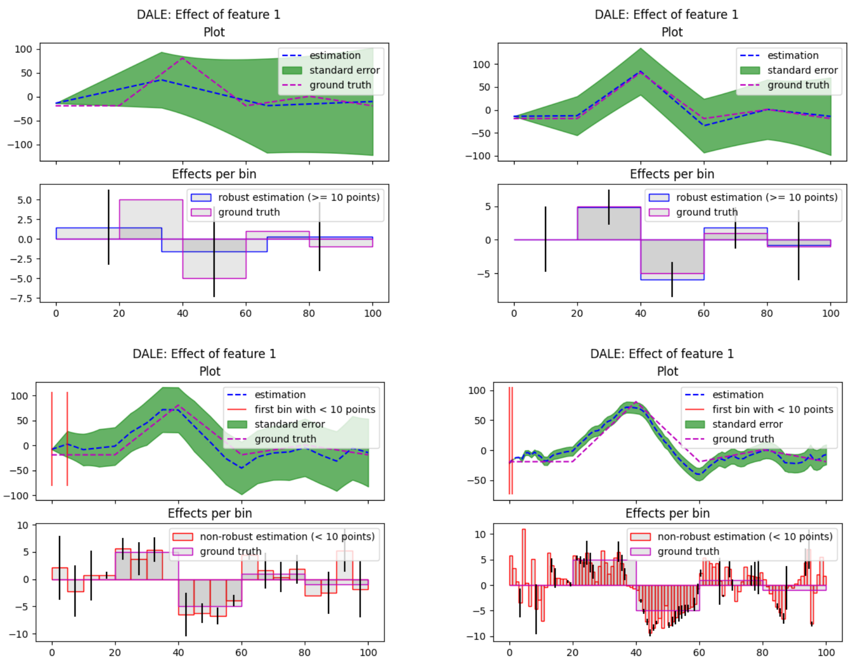

Figure 2: DALE effect with standard error for (a) $K = 3$, (b) $K = 5$, (c) $K = 20$, (d) $K = 100$

example, if we repeat the experiment with a bigger dataset ($N = 10000$) points, we get the results of figure 3. We observe that standard error gives accurate trustability regions, in all plots, i.e. (b), (c), (d), apart from (a). In (a), (b) and (c), it correctly us to trust the feature effect plots. Misleadingly, it also informs to trust plot (a), which is wrong. This because in this case, the first constraint has been violated.

Conclusion

If the user can create small bins, i.e. (dx small enough $\rightarrow K$ big enough, to respect constraint 1 and has enough points inside each bin (constraint 2), then the standard error reveals to what extend we can trust the feature effect plot.

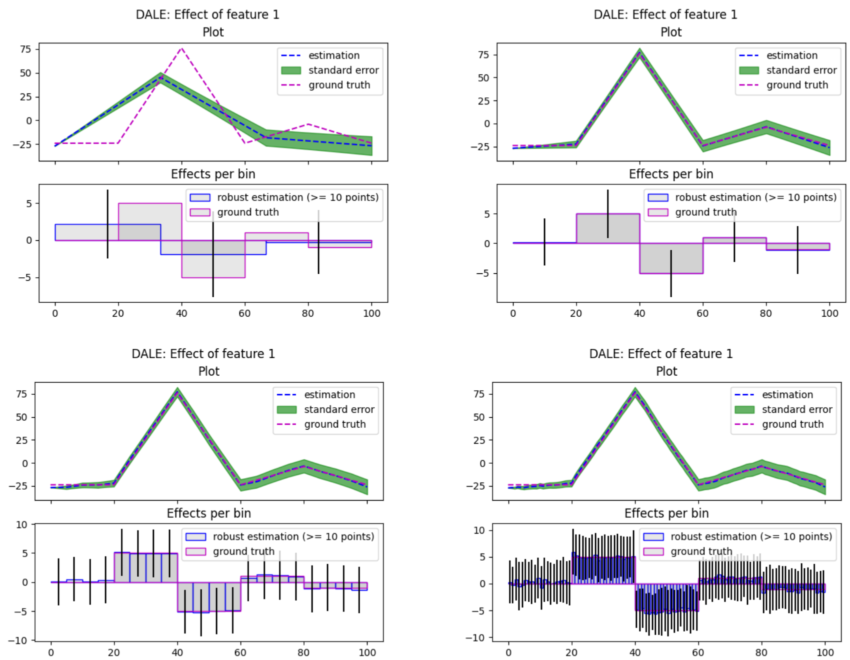

Figure 3: DALE effect with standard error for (a) $K = 3$, (b) $K = 5$, (c) $K = 20$, (d) $K = 100$

3 Choose the more accurate feature effect plot to trust

Statement

Apart from having a metric of uncertainty about the feature effect plot (standard error), we must also have a metric to choose the most accurate feature effect plot.

Proposal

We propose the minimization of the accumulated standard deviation (or accumulated variance). For ALE with K bins, let's notate as:

- dx^K the length of each bin
- p_i^K the number of the training points inside the $i-th$ bin
- σ_i^K the std of the local effects of the training points inside the $i-th$ bin

We will minimize:

$$K_{min} = \text{argmin}_K \quad [dx^K \sum_i^K \sigma_i^K * (1 - d_i^K)] \tag{4}$$

$$\text{s.t. } p_i^K \geq \texttt{min\_points\_per\_bin} \ \forall i \tag{5}$$

where $d_i^K = 0.2 * \frac{p_i^K}{N} \in [0, 0.1]$ works as a 'discount', favoring the creation of bigger bins in cases of similar standard deviation.

Example

Let's see how the proposed metric works in the example of Chapter 2. We set `min_points_per_bin = 10`. We repeat the procedure for different number of points $N = 50, 100, 1000, 10000$. The results are shown in 4. For $N = 50$ we can evaluate the loss only up to 4 bins (3 is the best) and for $N = 100$ up to 7 (5 is the best). In all other cases 5 bins is the best solution (sensible), with all multiples of 5 are also good options. All results make sense.

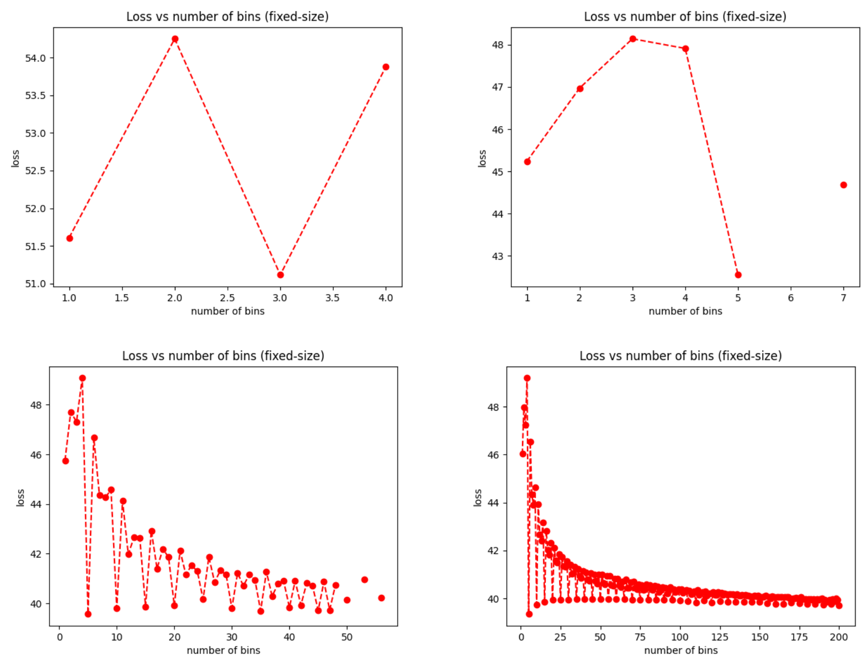

Figure 4: Loss (=accumulated standard error) for datasets of size (a) $N = 50$, (b) $N = 100$, (c) $N = 1000$, (d) $N = 10000$

Conclusion

Minimizing the accumulated standard deviation is a good indicator for choosing the most accurate feature effect plot.

4 Variable-size bins are important in many cases

Statement

Spliting the space in equally sized-bins is not always a good option.

Example

Let's see the example of Figure 5, where the ground truth effect is a piecewise linear function with 3 parts. There are two problems with applying equally-sized bins:

- The first two parts have length 20, whereas the third part has length 60. For capturing the first two effects we need $dx \leq 20$, but this resolution adds noise to the third bin that could be bigger.
- the data points are not uniformly split along the axis. From x=40 until x=100, they are sparse. This makes it difficult to split the space, with many bins since we won't have enough points per bin.

Therefore as we can see in figure 6, with a threshold of $min\_points\_per\_bin = 10$, we can evaluate until $K = 4$, which has the best loss. As we see, $K = 4$ has lower resolution than needed, badly estimating feature effect (Figure 6(b)).

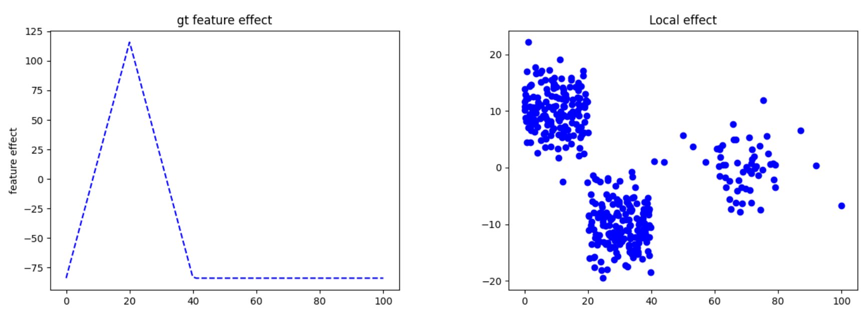

Figure 5: Left: Ground truth feature effect plot, Right: local effect for each point

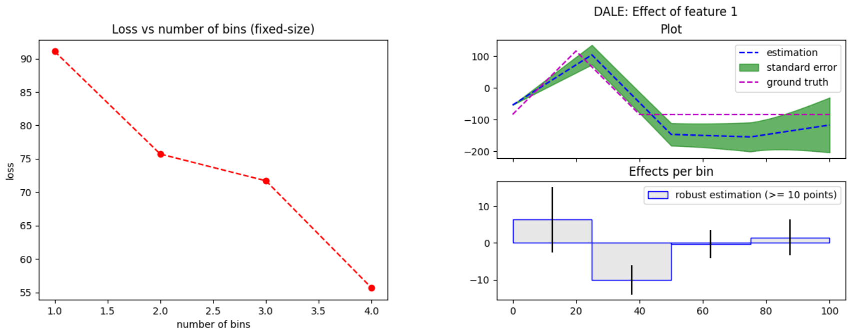

Figure 6: Left: Ground truth feature effect plot, Right: local effect for each point

Proposal

We propose the creation of variable-size bins. The objective is the minimisation of the loss as described in Section 3. In our example, the algorithm created 5 bins (figure 7 (a)) that lead to the feature effect of (figure 7 (b)), which is almost perfect.

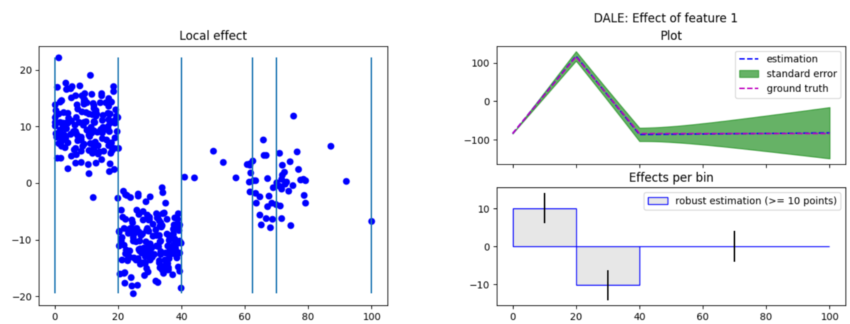

Figure 7: Left: Ground truth feature effect plot, Right: local effect for each point

Conclusion

5 Variable-size bins in case of non-linear feautre effects